

Policy Based Reinforcement Learning: A Tutorial


Peng Sun pythonsun@tencent.com

Vision Group, Tencent AI Lab

(Deep) Reinforcement Learning

Viewing (raw) states, making serial decisions

Playing Video Game



In-door Robot Navigation



Robot Arm Control




Policy Based (Deep) Reinforcement Learning

Learning the policy directly

A function approximator $\pi(a|s; \theta)$, mapping raw states s to actions a

E.g., the Atari game “breakout”

- Raw states s is 3x84x84 RGB image
- Actions {left, right, fire, no-op}



How It Works: a Quick Overview

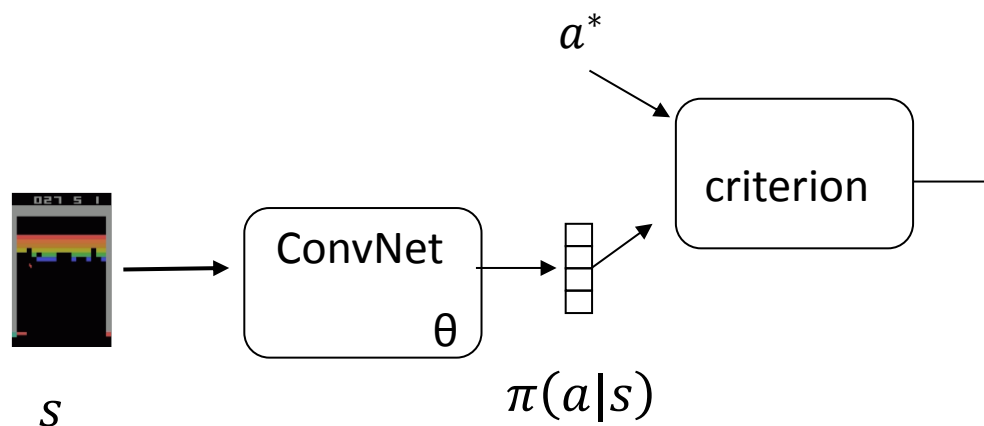
Inputs: image s

Outputs: policy probability $\pi(a|s; \theta)$

Ground Truth: action a^*

Criterion: class log likelihood $\log p(a^*|s; \theta)$

Gradients: $\nabla_{\theta} \log p(a^*|s; \theta)$



Supervised learning

How It Works: a Quick Overview (Cont.)

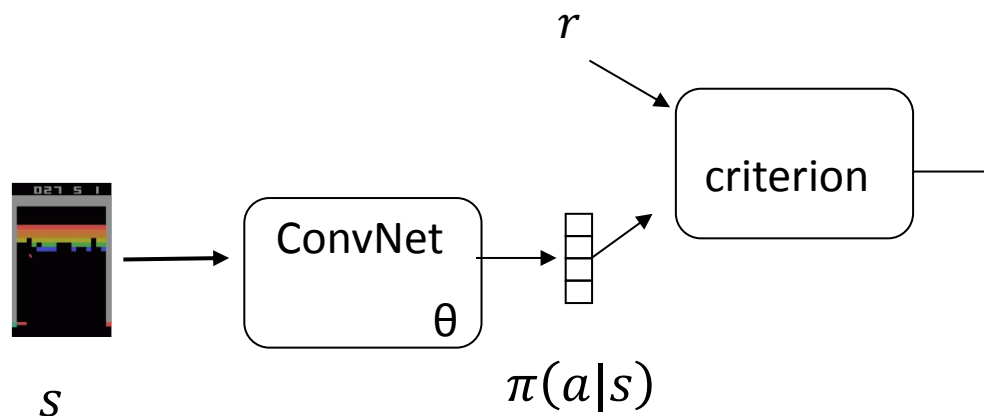
Inputs: image s

Outputs: policy probability $\pi(a|s; \theta)$

Ground Truth: reward r

Criterion: sample $a \sim \pi(a|s; \theta)$, then
weighted log likelihood $r \log p(a|s; \theta)$

Gradients: $\nabla_{\theta} r \log p(a|s; \theta)$



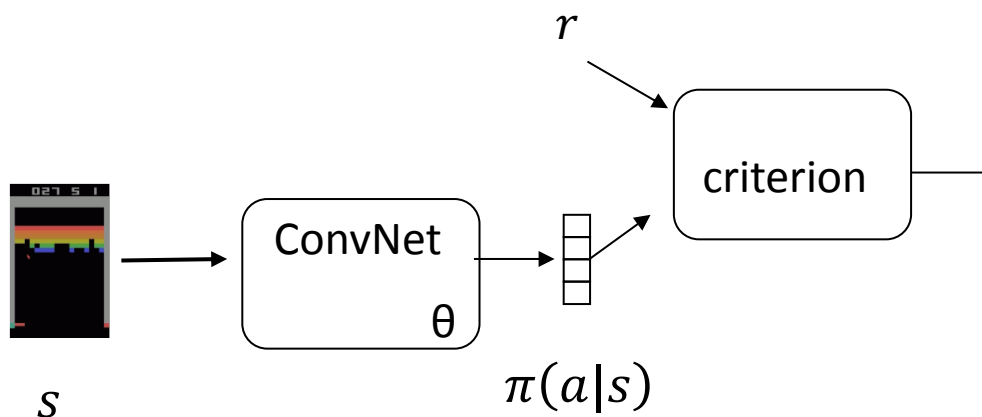
Reinforcement learning

How It Works: a Quick Overview (Cont.)

Gradients: $\nabla_{\theta} r \log p(a|s; \theta)$

- Estimation quality (bias, variance)?
- Delayed reward/credit assignment?

Need formalization




Reinforcement learning

Outline

- Log-Likelihood Objective Optimization
- “Vanilla” Policy Gradient
- Asynchronous Advantage Actor-Critic (A3C)
- Trust Region Policy Optimization (TRPO)
- Guided Policy Search

Agent Environment Interaction



	s_{t-1}	s_t	
\dots	a_{t-1}	a_t	\dots
	r_{t-1}	r_t	

Markovian Process


Assumptions:

current action over current state: $a_t \sim \pi(a_t | s_t)$

current state over previous state, action: $s_t \sim p(s_t | s_{t-1}, a_{t-1})$

current reward over current state, action: $r_t = r(s_t, a_t)$

The trajectory



Probabilistic Graphical Model


Arrows describe the dependency

Easy to write down how to decompose the joint distribution

Random variables: $x = (x_1, x_2, x_3, x_4)$

The joint distribution:

$$p(x_1, x_2, x_3, x_4) = p(x_1) p(x_2|x_1)p(x_3|x_2)p(x_4|x_2, x_3)$$




Learning

What to approximate?

- Function approximator of the policy $\pi(a|s; \theta)$

What to learn?

- Parameters θ
- Altering $\pi(a|s; \theta)$ changes the whole trajectory...




$$s_t \sim p(s_t | s_{t-1}, a_{t-1})$$

$$a_t \sim \pi(a_t | s_t)$$

$$r_t = r(s_t, a_t)$$

Learning Objective



$x = (s, a)$ plane

Define the objective

$$U(\theta) = \mathbb{E}_{\tau}[R(\tau)]$$

$\tau = (s_1, a_1, \dots)$ the trajectory

$R(\tau)$ the rewards over τ e.g., the cumulated rewards

The expectation: sums over all possible s_1, a_1, \dots

Derive the gradient

Objective

$$U(\theta) = \sum_{\tau} P(\tau; \theta) R(\tau)$$

Gradient

$$\begin{aligned}\nabla_{\theta} U(\theta) &= \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \nabla_{\theta} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \frac{P(\tau; \theta)}{P(\tau; \theta)} \nabla_{\theta} P(\tau; \theta) R(\tau)\end{aligned}$$

Population mean


$$\begin{aligned}&= \sum_{\tau} P(\tau; \theta) \frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)} R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)\end{aligned}$$

Sample mean


$$\nabla_{\theta} U(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) R(\tau^{(i)})$$

Decompose the gradient

$$\begin{aligned}\nabla_{\theta} \log P(\tau^{(i)}; \theta) &= \nabla_{\theta} \log \left[\prod_{t=0}^H \underbrace{P(s_{t+1}^{(i)} | s_t^{(i)}, u_t^{(i)})}_{\text{dynamics model}} \cdot \underbrace{\pi_{\theta}(u_t^{(i)} | s_t^{(i)})}_{\text{policy}} \right] \\&= \nabla_{\theta} \left[\sum_{t=0}^H \log P(s_{t+1}^{(i)} | s_t^{(i)}, u_t^{(i)}) + \sum_{t=0}^H \log \pi_{\theta}(u_t^{(i)} | s_t^{(i)}) \right] \\&= \nabla_{\theta} \sum_{t=0}^H \log \pi_{\theta}(u_t^{(i)} | s_t^{(i)}) \\&= \sum_{t=0}^H \underbrace{\nabla_{\theta} \log \pi_{\theta}(u_t^{(i)} | s_t^{(i)})}_{\text{no dynamics model required!!}}\end{aligned}$$



Remarks/Propositions



$$\nabla_{\theta} U(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) R(\tau^{(i)})$$

- make high reward trajectory more possible; low reward trajectory less possible
- $R(\tau)$ can be non-decomposable, or even non-continuous!!

Application: Image Captioning

Image captioning as a sequential decision making task

- non-differentiable metric CIDER

The trajectory

- State is the hidden cells of LSTM
- Action is what word to choose
- A trajectory corresponds to a sentence

$$\begin{array}{ccccc} & s_{t-1} & s_t & & \\ \dots & & & & \dots \\ & a_{t-1} & a_t & & \end{array}$$

Outline

- Log-Likelihood Objective Optimization
- “Vanilla” Policy Gradient
- Asynchronous Advantage Actor-Critic (A3C)
- Trust Region Policy Optimization (TRPO)
- Guided Policy Search

Bias and Variance

$$\nabla U(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) R(\tau^{(i)})$$

- The gradient estimate is unbiased.
- But it's noisy, i.e., high variance
- Develop techniques to lower variance and make it practical in real word application
 - Baseline
 - Temporal structure

Introduce baseline

$$\nabla U(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) R(\tau^{(i)})$$

- If R happens to be non-negative, all trajectories would be always pushed.
- How to push only the “good enough” trajectories?

$$\nabla U(\theta) \approx \hat{g} = \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) (R(\tau^{(i)}) - b)$$

- Push good enough trajectory (bigger than b), avoid poor trajectory (less than b)
- Much like why we need a bias term for binary linear classifier
- Reasonable $b = \mathbb{E}[R(\tau)] \approx \frac{1}{m} \sum_{i=1}^m R(\tau^{(i)})$

Consider temporal structure

Current policy be responsible for all rewards

$$\begin{aligned}\hat{g} &= \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^{(i)}; \theta) (R(\tau^{(i)}) - b) \\ &= \frac{1}{m} \sum_{i=1}^m \left(\sum_{t=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u_t^{(i)} | s_t^{(i)}) \right) \left(\sum_{t=0}^{H-1} R(s_t^{(i)}, u_t^{(i)}) - b \right)\end{aligned}$$

Intuitive: current policy be only responsible for future rewards, not the past awards

$$\frac{1}{m} \sum_{i=1}^m \sum_{t=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u_t^{(i)} | s_t^{(i)}) \left(\sum_{k=t}^{H-1} R(s_k^{(i)}, u_k^{(i)}) - b(s_k^{(i)}) \right)$$

- R term over the future
- Baseline term the expectation over the future

$$b(s_t) = \mathbb{E} [r_t + r_{t+1} + r_{t+2} + \dots + r_{H-1}]$$

Pseudo code for “Vanilla” Policy Gradient

Algorithm 1 “Vanilla” policy gradient algorithm

Initialize policy parameter θ , baseline b

for iteration=1, 2, ... **do**

Collect a set of trajectories by executing the current policy

At each timestep in each trajectory, compute

the *return* $R_t = \sum_{t'=t}^{T-1} \gamma^{t'-t} r_{t'}$, and

the *advantage estimate* $\hat{A}_t = R_t - b(s_t)$.

Re-fit the baseline, by minimizing $\|b(s_t) - R_t\|^2$,
summed over all trajectories and timesteps.

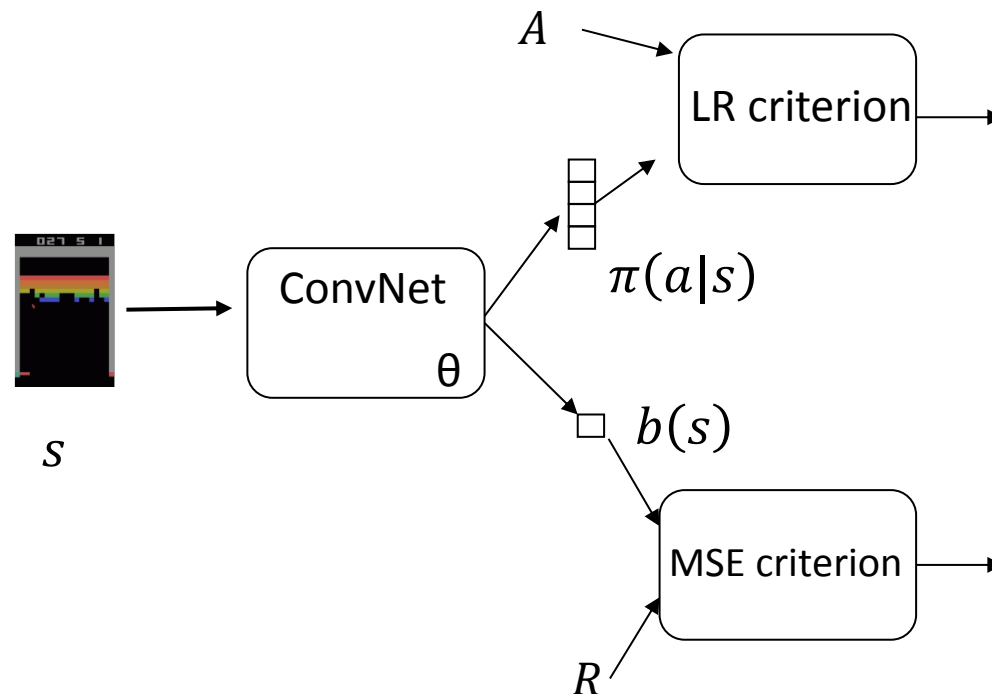
Update the policy, using a policy gradient estimate \hat{g} ,
which is a sum of terms $\nabla_{\theta} \log \pi(a_t | s_t, \theta) \hat{A}_t$

end for

Remarks: actually works very well... will see this next...

Neural Network Implementation

The “pseudo ground truth” A and R must be generated on-the-fly.



Outline

- Log-Likelihood Objective Optimization
- “Vanilla” Policy Gradient
- **Asynchronous Advantage Actor-Critic (A3C)**
- Trust Region Policy Optimization (TRPO)
- Guided Policy Search

Advantage Actor Critic

Almost the same with “vanilla” policy gradient, but different terminology.

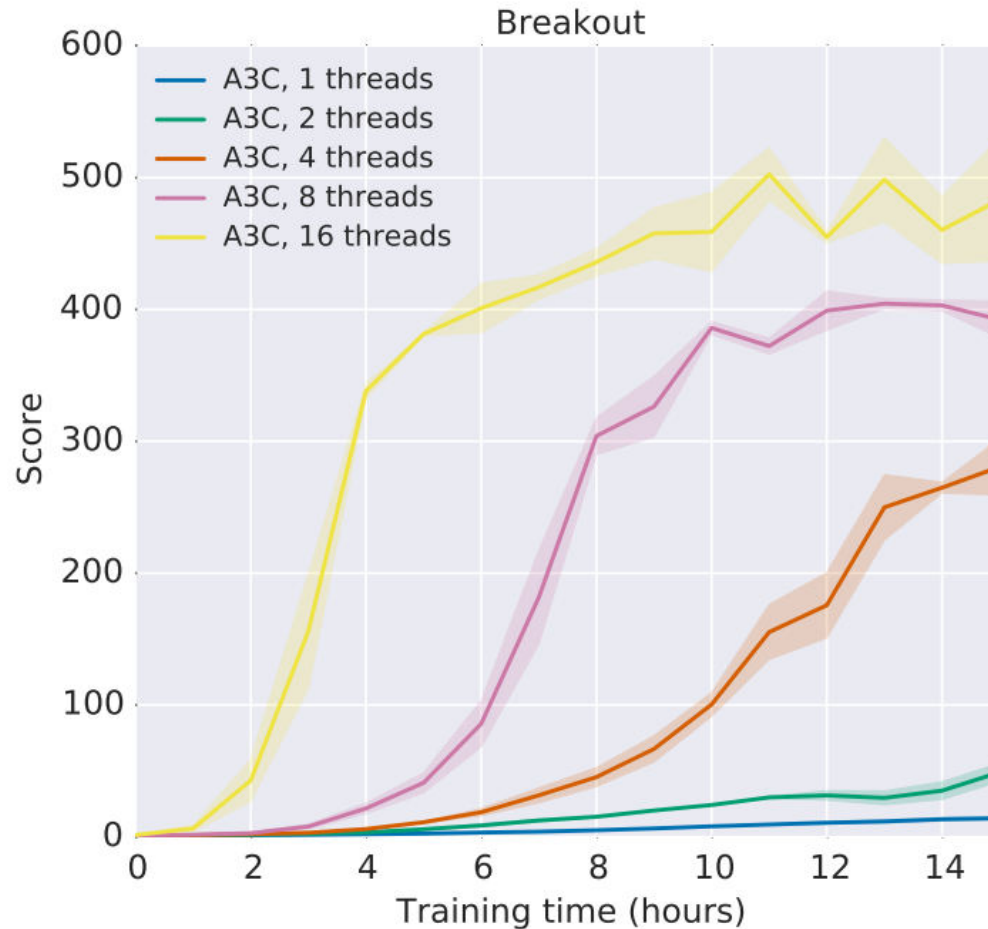
- Actor: the policy $\pi(a|s; \theta)$
- Critic: the baseline $V(s)$, criticize how the state is
- Advantage: the minus term

Asynchronous Advantage Actor Critic (A3C)

- Start many agent-environment interaction threads
- Only one nn, whose parameters are shared across threads
- Parameters are updated asynchronously during training
 - easier implementation, as it is lock-free
 - unreasonable in regards of numeric accuracy, but work well in practice.


Asynchronous Advantage Actor Critic (A3C): performance

Significant speed-up (figure)



Demo: Atari Breakout

Show movie here



Demo: 3D car racing

<https://youtu.be/0xo1Ldx3L5Q>

Remark: the same nn and RL algo with the 2D Atari Game!

Demo: ViZDoom



Single-thread also works well, although much slower

- This is not true for other RL method (e.g., Q-learning), which requires high-quality random sampling. Multi threading ensures this.

Strong empirical results

- A3C is uniformly better than other Q-learning based methods
- Could be due to the introduction of $V(s)$ term.
- Didn't compare to asynchronous dueling-network using similar idea

Anyway, A3C should be an off-the-shelf RL method for problem on hand! Why it was “overlooked” before, though? Hmm...

Outline

- Log-Likelihood Objective Optimization
- “Vanilla” Policy Gradient
- Asynchronous Advantage Actor-Critic (A3C)
- Trust Region Policy Optimization (TRPO)
- Guided Policy Search

Step Size Matters

Too small step size: very slow

Too big step size:

- Supervised Learning: fixed at next step
- Reinforcement Learning: lead to poor state regions; sample poor trajectories; cannot fix...



Trust Region Policy Optimization

During each iteration

- Get trajectories $\{s_n, a_n, r_n\}$ using $\pi(\cdot | \cdot; \theta_{old})$
- Maximize

$$L(\theta),$$
$$\text{s.t. } D(\theta_{old}, \theta) < \delta$$

- Update $\theta_{old} \leftarrow \theta$

The sample version, surrogate “local” rewards

$$L(\theta) = \sum_{n=1}^N \frac{\pi(a_n | s_n; \theta)}{\pi(a_n | s_n; \theta_{old})} \widehat{A}_n$$

The sample version, average constraint

$$D(\theta_{old}, \theta) = \sum_{n=1}^N KL(\pi(\cdot | s_n; \theta_{old}), \pi(\cdot | s_n; \theta))$$

The numeric solver

Objective $L(\theta)$ expanded to first order, get its gradient g

Constraint $D(\theta_{old}, \theta)$

- KL: zero-order and first-order expansion are
- Second order expansion, get Hessian matrix F

Find the updating direction $s = F^{-1}g$

Do line search along s by $\theta \leftarrow \theta + \eta s$, step size η ensuring the constraint on $D(\theta_{old}, \theta)$ and improvement over $L(\theta)$

- Similar to Natural Gradient (NG): updating by $\theta \leftarrow \theta + \lambda F^{-1}g$
- TRPO turns out to outperform NG (although the subtle difference!).
- Careful step size is critical!

Demo

URL here

Outline

- Log-Likelihood Objective Optimization
- “Vanilla” Policy Gradient
- Asynchronous Advantage Actor-Critic (A3C)
- Trust Region Policy Optimization (TRPO)
- Guided Policy Search

Background

Robot control

- States: positions, velocity, joint angles...
- Actions: torques
- Rewards: negative distance to expected positions, angles (by design)

RL training on real environment, instead of simulator. Low “put-through” due to physics, even though the policy gradient algorithm itself can work very fast.

Require data-efficient training!

Data Efficiency v.s. Domain Knowledge


Why PG consumes so many training data?

- Does not model the environment dynamics (model-free).

Explicitly model the environment dynamics

- Can hopefully learn with fewer data
- E.g., parabolic curve fitting

Optimal Control/Trajectory Optimization



Given an initial state, find a good trajectory along which high rewards are collected.

- Simultaneous rollout and local policy $q_i(a_t|s_t)$ learning
- Analytic form! E.g., Dynamics is Gaussian, reward is quadratic, then $q_i(a_t|s_t)$ is time-varying linear Gaussian
- $q_i(a_t|s_t)$ good for current trajectory, but generalizes poorly to other states or trajectories

General Pipeline

Given

- Known environment dynamics $p(a_t | s_{t-1}, a_{t-1})$ or fitted from data
- Known reward function $r_t = r(s_t, a_t)$ by design

When not convergent, iterate over

1. Sample on initial state and do trajectory optimization to get trajectories and $q_1(a|s), q_2(a|s), \dots$
2. Supervised Learning: Fit $\pi(a|s; \theta)$ to $q_1(a|s), q_2(a|s), \dots$ on the points on the trajectories

The local policy is called guided policy, hence the name

Demo

<https://youtu.be/mSzEyKaJTSU>

Visual Motor Control

Control robot by viewing image sequence

Methodology: Guided Policy Search with asymmetrical states in the two phases

- Guiding policy $q_i(a|s)$ over “intrinsic” states (position, velocity...)
- Learn policy $\pi(a|s; \theta)$ over “visual” states (pixels from camera)

Remarks

- A promising framework to exploit environment dynamics
- How to generalize to other domains than robotics? (E.g., playing 3D video game)